

Maven

```
<dependency>
  <groupId>org.mapdb</groupId>
  <artifactId>mapdb</artifactId>
  <version>[version]</version>
</dependency>
```

Map stored in file

```
import org.mapdb.*;
```

```
//open (or create) database
File file = new File("dbFileName");
DB db = DBMaker
    .newFileDB(file)
    .make();
```

```
//use map
Map map = db.getHashMap("mapName");
map.put("aa","bb");
```

```
//commit and close database
db.commit();
db.close();
```

In-memory off-heap Map

```
// same as above except different method
DB db = DBMaker
    .newMemoryDirectDB();
    .make();
```

In-memory off-heap Queue

```
// same as above except different method
DB db = DBMaker
    .newMemoryDirectDB();
    .make();
Queue q = db.getQueue("q");
```

Options to make it faster

```
DB db = DBMaker
// all options works with files as well
    .newMemoryDB();
// disable transactions make writes
// but you may lose data if store crashes
    .transactionsDisable()
// memory mapped files are faster
// but does not work well on 32bit
    .mmapFileEnable()
// writes done by background thread
// it has some overhead, so could be slower
    .asyncWriteEnable()
// increase cache size if you have free heap
// default value is 32K
    .cacheSize(1000000)
    .make();
```

Other DBMaker options

```
// encrypt data with password
    .encryptionEnable("password")
// use fast compression
    .compressionEnable()
// enables CRC32 checksum
// to protect from data corruption
    .checksumEnable()
```

Cache options

```
// It caches deserialized objects on heap.
// Default cache size is 32,000, increase it
    .cacheSize(1000000)
// enable least-recently-used cache
    .cacheLRUEnable()
// Unbounded soft-reference cache
// use with plenty of free heap
    .cacheSoftRefEnable()// Hard ref, use if
// heap is larger then store
    .cacheHardRefEnable()
```

Concurrent transactions

```
// By default there is single-global
// transaction per store.
// This enables proper transactions
// with full serializable isolation
TxMaker txMaker = DBMaker
    .newFileDB(file)
    .makeTxMaker();

// open two transactions, with single map
// both can only see their own changes
DB tx1 = txMaker.makeTx();
Map map1 = tx1.getTreeMap("map");
DB tx2 = txMaker.makeTx();
Map map2 = tx2.getTreeMap("map");

//commit and close
tx1.commit()
tx2.commit()
txMaker.close()
```

Snapshots

```
// lighter way to get consistent data view
DB db = DBMaker
    .newFileDB(file)
    .snapshotEnable()
    .make()
Map map = db.getHashMap("map");
map.put(1,2);
DB snap = db.snapshot();

Map mapOld = snap.getHashMap("map");
map.put(3,4); //mapOld still has only 1,2
snap.close(); //release resources
```

```
// Third way to ensure consistency is
// Compare and Swap operation. MapDB
// has ConcurrentMap and atomic variables.
```

Maps and Sets

```
// Shows how to get all available collections
DB db = DBMaker
    .newMemoryDirectDB();
    .make();
```

```
// BTreeMap is good for small sorted keys
ConcurrentNavigableMap treeMap =
    db.getTreeMap("treeMap");
```

```
// HashMap (aka HTreeMap) is good for
// larger keys and as a cache
ConcurrentMap hashMap =
    db.getHashMap("hashMap");
```

```
// there is also TreeSet and HashSet
SortedSet treeSet = db.getTreeSet("ts");
Set hashSet = db.getHashSet("hashSet");
```

Queues

```
// first-in-first-out queue
BlockingQueue fifo = db.getQueue("fifo");
```

```
// last-in-first-out queue (stack)
BlockingQueue lifo = db.getStack("lifo");
```

```
// circular queue with limited size
BlockingQueue c =
    db.getCircularQueue("circular");
```

Atomic records

```
// atomically updated records stored in DB
// Useful for example for sequential IDs.
// there is Long, Integer, String
// and general atomic variable
Atomic.Long q = db.getAtomicLong("long");
q.set(1999);
long id = q.incremendAndGet();
```

Configuring maps

```
// create map optimized for large values
Map<String,String> m =
    db.createTreeMap("treeMap");

//serializers are critical for performance
    .keySerializer(BTreeKeySerializer.STRING)
// compress large ASCII string values
    .valueSerializer(
        new Serializer.CompressionWrapper(
            Serializer.STRING_ASCII))
// and store values outside of BTree nodes
    .valuesOutsideNodesEnable()
// enable size counter
    .counterEnable()
// make BTree nodes larger
    .nodeSize(120)
// and finally create map
    .makeOrGet();
```

Secondary indexes

```
// create secondary value (1:1 relation)
// secondary map gets auto updated
Map<ID, Person> persons
Map<ID, Branch> branches
Bind.secondaryValue(persons,branches,
    (person)-> person.getBranch());

// create secondary key (index) for age(N:1)
SortedSet<Fun.Tuple2<Age,ID>> ages
Bind.secondaryKey(persons, ages,
    (person)-> person.getAge());

// get all persons of age 32
for(ID id: Fun.filter(ages, 32)){
    Person p = persons.get(id)
}
```

HTreeMap as a cache

```
// Entries are removed if map is too large
```

```
// Off-heap map with max size 16GB
Map cache = DBMaker
    .newCacheDirect(16)
```

```
// On-disk cache in temp folder
// with max size 128GB or 1M entries
```

```
DB db = DBMaker
    .newTempFileDB()
    .transactionDisable()
    .closeOnJvmShutdown()
    .deleteFilesAfterClose()
    .make()
Map cache = db
    .createHashMap("cache")
    .expireStoreSize(128) // GB
    .expireMaxSize(1000000)
    .make()
```

Data Pump for faster import

```
// Data Pump creates TreeMap and TreeSet
// in streaming fashion. Import time is linear
// to number of entries.
```

```
Iterator iter = ... iterate over keys..
```

```
Map<K,V> m = db.createTreeMap("map")
    .pumpSource(iter, (key)-> key.getValue())
    .pumpIgnoreDuplicates()
    .pumpPresort(1000000)
    .make()
```